

---

# Overview

---

- ❖ Types of statistical learning problems
- ❖ Why learning is difficult?
- ❖ Bias variance tradeoff
- ❖ An example: kNN *vs* Linear Regression (in a separate pdf file)
- ❖ Not all about prediction

---



# Bias Variance Tradeoff

---

Goal of ML: Minimize *generalization error* (i.e., error on unseen future datasets), not training error.

Source of errors:

- ❖ Bias
- ❖ Variance




# Bias Variance Tradeoff

Goal of ML: Minimize *generalization error* (i.e., error on unseen future datasets), not training error.

Source of errors:

- ❖ Bias
- ❖ Variance




# Bias Variance Tradeoff

Goal of ML: Minimize *generalization error* (i.e., error on unseen future datasets), not training error.

Source of errors:

- ❖ Bias
- ❖ Variance




# Bias Variance Tradeoff

Goal of ML: Minimize *generalization error* (i.e., error on unseen future datasets), not training error.

Source of errors:

- ❖ Bias
- ❖ Variance




# Bias Variance Tradeoff

Goal of ML: Minimize *generalization error* (i.e., error on unseen future datasets), not training error.

Source of errors:

- ❖ Bias
- ❖ Variance



---

# What'll be Covered in Stat542


---

- ❖ Flexible modeling techniques to reduce bias
- ❖ Useful strategies to achieve the tradeoff between bias and variance

# Two Successful Strategies

---


- ❖ Regularization: Restrict the parameters to a low-dimensional space, which is *adaptively* determined by the data.





# Two Successful Strategies

- ❖ Regularization: Restrict the parameters to a low-dimensional space, which is *adaptively* determined by the data.




$$\min_w \left[ \sum_{i=1}^n \text{loss}(y_i, f(x_i; w)) + \lambda |w| \right]$$

LASSO

# Two Successful Strategies

- ❖ Regularization: Restrict the parameters to a low-dimensional space, which is *adaptively* determined by the data.
- ❖ Ensemble: Average many low-bias high-variance models; averaging reduces variance.



**FIGURE 2.11.** Test and training error as a function of model complexity.

---

# Overview

---

- ❖ Types of statistical learning problems
- ❖ Why learning is difficult?
- ❖ Bias variance tradeoff
- ❖ An example: kNN *vs* Linear Regression
- ❖ Not all about prediction

---

# Not All About Prediction

---

- ❖ Although the focus of this course is prediction, statistical learning  $\neq$  prediction
- ❖ Exploration *vs.* Prediction
- ❖ Data product *vs.* decision making
- ❖ Make your model to generate actionable insights